

Navigation

June 23, 2021

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip -q install ./python
```

```
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatible.  
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 3.0.
```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```
In [9]: from unityagents import UnityEnvironment  
import numpy as np
```

```
# please do not modify the line below
```

```
env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")
```

```
INFO:unityagents:
```

```
'Academy' started successfully!
```

```
Unity Academy name: Academy
```

```
Number of Brains: 1
```

```
Number of External Brains : 1
```

```
Lesson number : 0
```

```
Reset Parameters :
```

```
Unity brain name: BananaBrain
```

```
Number of Visual Observations (per agent): 0
```

```
Vector Observation space type: continuous
```

```
Vector Observation space size (per agent): 37
```

```

Number of stacked Vector Observation: 1
Vector Action space type: discrete
Vector Action space size (per agent): 4
Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [10]: # get the default brain
         brain_name = env.brain_names[0]
         brain      = env.brains[brain_name]

```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [11]: # reset the environment
         env_info = env.reset(train_mode=True)[brain_name]

         # number of agents in the environment
         print('Number of agents:', len(env_info.agents))

         # number of actions
         action_size = brain.vector_action_space_size
         print('Number of actions:', action_size)

         # examine the state space
         state = env_info.vector_observations[0]
         print('States look like:', state)
         state_size = len(state)
         print('States have length:', state_size)

```

```

Number of agents: 1
Number of actions: 4
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755     1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
States have length: 37

```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that in this coding environment, you will not be able to watch the agent while it is training, and you should set `train_mode=True` to restart the environment.

```
In [12]: env_info = env.reset(train_mode=True)[brain_name] # reset the environment
state = env_info.vector_observations[0] # get the current state
score = 0 # initialize the score
while True:
    action = np.random.randint(action_size) # select an action
    env_info = env.step(action)[brain_name] # send the action to the environment
    next_state = env_info.vector_observations[0] # get the next state
    reward = env_info.rewards[0] # get the reward
    done = env_info.local_done[0] # see if episode has finished
    score += reward # update the score
    state = next_state # roll over the state to next time
    if done: # exit loop if episode finished
        break

print("Score: {}".format(score))
```

Score: 0.0

When finished, you can close the environment.

```
In [23]: env.close()
```

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

```
In [1]: '''
        Note to the reviewer 23-Jun-21

        This project used a double deep-Q network modifying the code provided in the lessons.
        Alterations to the code were in the double_dqn, the agent, and the class as per the file
        submitted. I used Batch Normalization to accelerate learning on layers one and two, altered
        the fcx_unit sizes which did very little to improve performance so I stuck with 64 as the
        fc1 and fc2. The execution was run in 'CPU' mode.
```

The final saved trained model used a three layer neural network where the batchsize was
Hyper-parameters:

(1) LR-learning-rate: LR = 5e-55
(2) Epsilon-Decay: eps_decay = 0.9645
(3) Gamma: GAMMA = 0.95
(4) Batch size: BATCH_SIZE = 32

Agent Parameters:

(1) state size: state_size = 37
(2) Env actions: action_size = 4
(3) Seed: seed = 42

scores = 0

These parameters realized the best results. Note: Dropping Gamma to around 0.95 helped re
faster, but not as much as adjusting Epsilon-Decay, i.e., slowing it and LR, i.e., speed

Within this notebook you will also see runs using a four layer neural net altering the
hyper-parameters. Results were roughly the same so I chose the final three layer model a
the following output:

```
BATCH_SIZE 32
Episode 100            Average Score: 2.33
Episode 200            Average Score: 7.11
Episode 300            Average Score: 9.86
Episode 393            Average Score: 13.07
Environment solved in 293 episodes!            Average Score: 13.07
```

Thank you.

Kim

'''

```
Out[1]: '\nNote to the reviewer 18-Jun-21\nThis project used a double deep-Q network modifying t
```

```
In [2]: #
# . Import the Necessary Packages
#
#
import random
import torch
import numpy as np
from collections import deque
import matplotlib.pyplot as plt
%matplotlib inline

!python -m pip install pyvirtualdisplay
from pyvirtualdisplay import Display
display = Display(visible=0, size=(1400, 900))
```

```
display.start()

is_ipython = 'inline' in plt.get_backend()
if is_ipython:
    from IPython import display

plt.ion()
```

Collecting pyvirtualdisplay

Downloading <https://files.pythonhosted.org/packages/79/30/e99e0c480a858410757e7516958e149285ea>

Collecting EasyProcess (from pyvirtualdisplay)

Downloading <https://files.pythonhosted.org/packages/48/3c/75573613641c90c6d094059ac28adb748560>

Installing collected packages: EasyProcess, pyvirtualdisplay

Successfully installed EasyProcess-0.3 pyvirtualdisplay-2.2

```
In [5]: print(torch.__version__)
```

0.4.0

```
In [7]: #
        # Re the Environment and Agent
        #
        env_info = env.reset(train_mode=True)[brain_name]

        # number of agents in the environment
        print('Number of agents:', len(env_info.agents))

        # number of actions
        action_size = brain.vector_action_space_size
        print('Number of actions:', action_size)

        # examine the state space
        state = env_info.vector_observations[0]
        print('States look like:', state)
        state_size = len(state)
        print('States have length:', state_size)
```

Number of agents: 1

Number of actions: 4

```
States look like: [ 0.          0.          1.          0.          0.16101955  1.          0.
 0.          0.          0.04571758  1.          0.          0.          0.
 0.2937662  0.          0.          1.          0.          0.14386636
 0.          0.          1.          0.          0.16776823  1.          0.
 0.          0.          0.04420976  1.          0.          0.          0.
 0.05423063  0.          0.          ]
```

States have length: 37

```

In [14]: #
         # Establish path to files
         #
         import os
         dir = os.getcwd()
         print(dir)

         from pathlib import Path #/home/workspace/ etc
         print(*Path(dir).iterdir(), sep="\n") # print files in the directory
         #
         # set the trained file name
         #
         file_name = dir + '/trained_model.pt'
         file_name

/home/workspace
/home/workspace/dqn_agent_prj1.py
/home/workspace/model_prj1.py
/home/workspace/python
/home/workspace/__pycache__
/home/workspace/unity-environment.log
/home/workspace/checkpoint.pth
/home/workspace/.ipynb_checkpoints
/home/workspace/Navigation.ipynb

Out[14]: '/home/workspace/trained_model.pt'

In [15]: def get_ns_rewrd_done(env_info):
         next_state = env_info.vector_observations[0] # get the next state
         reward      = env_info.rewards[0]           # get the reward
         done        = env_info.local_done[0]        # get done-status
         return next_state, reward, done

In [16]: #nxt,rr,dd = get_ns_rewrd_done(env_info)
         nxt,_,_ = get_ns_rewrd_done(env_info)
         print(nxt,len(nxt))
         #print(rr)
         #print(dd)

[ 1.          0.          0.          0.          0.43962687  1.          0.
  0.          0.          0.19398789  1.          0.          0.          0.
  0.48112735  0.          0.          1.          0.          0.52109712
  0.          0.          1.          0.          0.38285938  1.          0.
  0.          0.          0.10405888  1.          0.          0.          0.
  0.37148571  0.          0.          ] 37

In [17]: from dqn_agent_prj1 import Agent
         agent = Agent(state_size=37, action_size=4, seed=42)

```

```

In [18]: def double_dqn(n_episodes=2000, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.99)
        """Deep Q-Learning.

        Params
        =====
        n_episodes (int): maximum number of training episodes
        max_t (int): maximum number of timesteps per episode
        eps_start (float): starting value of epsilon, for epsilon-greedy action selection
        eps_end (float): minimum value of epsilon
        eps_decay (float): multiplicative factor (per episode) for decreasing epsilon
        """
        scores = [] # list containing scores from each episode
        scores_window = deque(maxlen=100) # last 100 scores
        eps = eps_start # initialize epsilon

        for i_episode in range(1, n_episodes+1):
            #state = env.reset()
            env_info = env.reset(train_mode=True)[brain_name] # reset the environment
            state, _, _ = get_ns_rewr_done(env_info) # get the current state
            score = 0

            for t in range(max_t):
                action = agent.act(state, eps) # Get action(s) for current state
                env_info = env.step(action)[brain_name] # Take action w/in the environment
                next_state, reward, done = get_ns_rewr_done(env_info) # get next state & r

                agent.step(state, action, reward, next_state, done)
                state = next_state
                score += reward
                if done:
                    break
            scores_window.append(score) # save most recent score
            scores.append(score) # save most recent score
            eps = max(eps_end, eps_decay*eps) # decrease epsilon
            print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_wi
            if i_episode % 100 == 0:
                print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(score
            if np.mean(scores_window)>13.0:
                print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.forma
                #torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
                torch.save(agent.qnetwork_local.state_dict(), file_name)
                break
        return scores

In [4]: #
        # Set directory for saving model
        #

In [11]: ****KEEP****


```

```

# New 4-layer model 10-Jun-2021
#
BATCH_SIZE = 64
LR          = 5e-55
GAMMA       = 0.95
agent       = Agent(state_size=37, action_size=4, seed=42)
scores      = 0
scores = dqn(n_episodes=500, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.9865)
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

Episode 100	Average Score: 2.57
Episode 200	Average Score: 7.95
Episode 300	Average Score: 11.82
Episode 400	Average Score: 14.28
Episode 500	Average Score: 16.38



```

In [21]: ****KEEP
# New 4-layer model 10-Jun-2021 3:10pm
#

```




```

# esp_decay = 0.9845; GAMMA 0.05; batchsize = 32
#
BATCH_SIZE = 32
LR          = 5e-55
GAMMA       = 0.95
agent       = Agent(state_size=37, action_size=4, seed=42)
scores      = 0
scores = dqn(n_episodes=500, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=0.9845)
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

Episode 100	Average Score: 2.55
Episode 200	Average Score: 8.72
Episode 300	Average Score: 11.83
Episode 400	Average Score: 15.03
Episode 500	Average Score: 16.55



```

In [22]: ****BEST RUN ***
# 13-Jun-2021 3-layer model 9:20 AM
#

```

```

# esp_decay = 0.9645; GAMMA 0.95; batchsize = 32
#
BATCH_SIZE = 32
LR          = 5e-55
GAMMA       = 0.95
agent       = Agent(state_size=37, action_size=4, seed=42)
scores      = 0
print('BATCH_SIZE', BATCH_SIZE)
scores = double_dqn(n_episodes=700, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

BATCH_SIZE 32


Episode 100 Average Score: 2.33

Episode 200 Average Score: 7.11

Episode 300 Average Score: 9.86

Episode 393 Average Score: 13.07

Environment solved in 293 episodes! Average Score: 13.07



In [13]: #

13-Jun-2021 3-layer model 9:20 AM

```

#
#  esp_decay =0.9845;  GAMMA 0.05; batchsize =32
#
#
BATCH_SIZE = 64
LR          = 5e-55
GAMMA       = 0.95
agent       = Agent(state_size=37, action_size=4, seed=42)
scores      = 0
print('BATCH_SIZE',BATCH_SIZE )
scores = double_dqn(n_episodes=700, max_t=1000, eps_start=1.0, eps_end=0.01, eps_decay=
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

BATCH_SIZE 64

Episode 100 Average Score: 0.38


Episode 200 Average Score: 5.22

Episode 300 Average Score: 9.63

Episode 400 Average Score: 12.16

Episode 474 Average Score: 13.03

Environment solved in 374 episodes! Average Score: 13.03



```
In [25]: env.close()
```

```
In [24]: #
         # check that model is saved
         #
         print(*Path(dir).iterdir(), sep="\n")#"/home/workspace/dog_project/images/"

/home/workspace/dqn_agent_prj1.py
/home/workspace/model_prj1.py
/home/workspace/python
/home/workspace/__pycache__
/home/workspace/unity-environment.log
/home/workspace/checkpoint.pth
/home/workspace/.ipynb_checkpoints
/home/workspace/Navigation.ipynb
/home/workspace/trained_model.pt
```

```
In [3]: #
        #    Improving Performance
        #
        '''
        To improve performance I would change the architecture back to a 4-layer neural network
        Batch Normalization to accelerates learning on just two layers one and three; see the tou
        results using a four-layer network. I would altered the fc3_unit size to 32 versus 64 a
        reapply these same hyperparameters:
            BATCH_SIZE = 32 #
            LR          = 5e-65 # Speed up learning a little more
            GAMMA       = 0.95
            agent       = Agent(state_size=37, action_size=4, seed=42)
        I would also try the memory prioritization algorithm to determine if there is an increas
        attainment.
        '''
```

```
Out[3]: '\nTo improve performance I would change the architecture back to a 4-layer neural network'
```

```
In [ ]:
```